



Intermediate to Zeek

NSF Cybersecurity Summit Training

Oct 23th, 2023

Who Are We?



Fatema Bannat Wala

- Security Engineer @ ESnet/LBNL
- 8 yrs of Network defense
- Zeek LT Member / Zeek Training Subgroup Lead

Christian Kreibich

- Zeek Technical Lead and LT member
- Principal Engineer at Corelight
- 20+ years of network security & research

Agenda



- What is a Zeek Cluster? - Fatema (15 mins)
- Become familiar with Cluster Configs - Fatema (45 mins)
- Zeek's Filesystem Layout - Fatema (30 mins)
<break 30 > 10:30am - 11am
- Running Zeek in a Cluster - Fatema (30mins)
- Customizing Zeek - Fatema (30mins)
- Zeek Package Manager (zkg) - Fatema (30mins)
<Lunch 1hr> 1pm - 2pm
- Logging Framework - Christian (30 mins)
- Notice Framework - Christian (30 mins)
- Telemetry Framework - Christian (30 mins)
<break 30> 3:30pm - 4pm
- New Zeek Management Framework - Christian (45 mins)
- Adjacent Technologies - Christian (15 mins)



Housekeeping

- Zeek's Official Documentation:
 - <https://docs.zeek.org/en/master/>
- Join Zeek Slack channel #training



Docker / Zeek Usage

- Docker How to doc
 - <https://github.com/zeek/zeek-training/tree/master/Intro-to-Zeek22>
- Docker image
 - Zeek v6.0.1 installed on the image-
<https://hub.docker.com/repository/docker/zeek/training>
 - # docker pull zeek/training:intro23
 - # docker run -it --privileged zeek/training:intro23
 - Path inside the container to training resources used during this training:
[/zeek/training-res/](#)
- Verify that Zeek runs correctly:
zeek -h

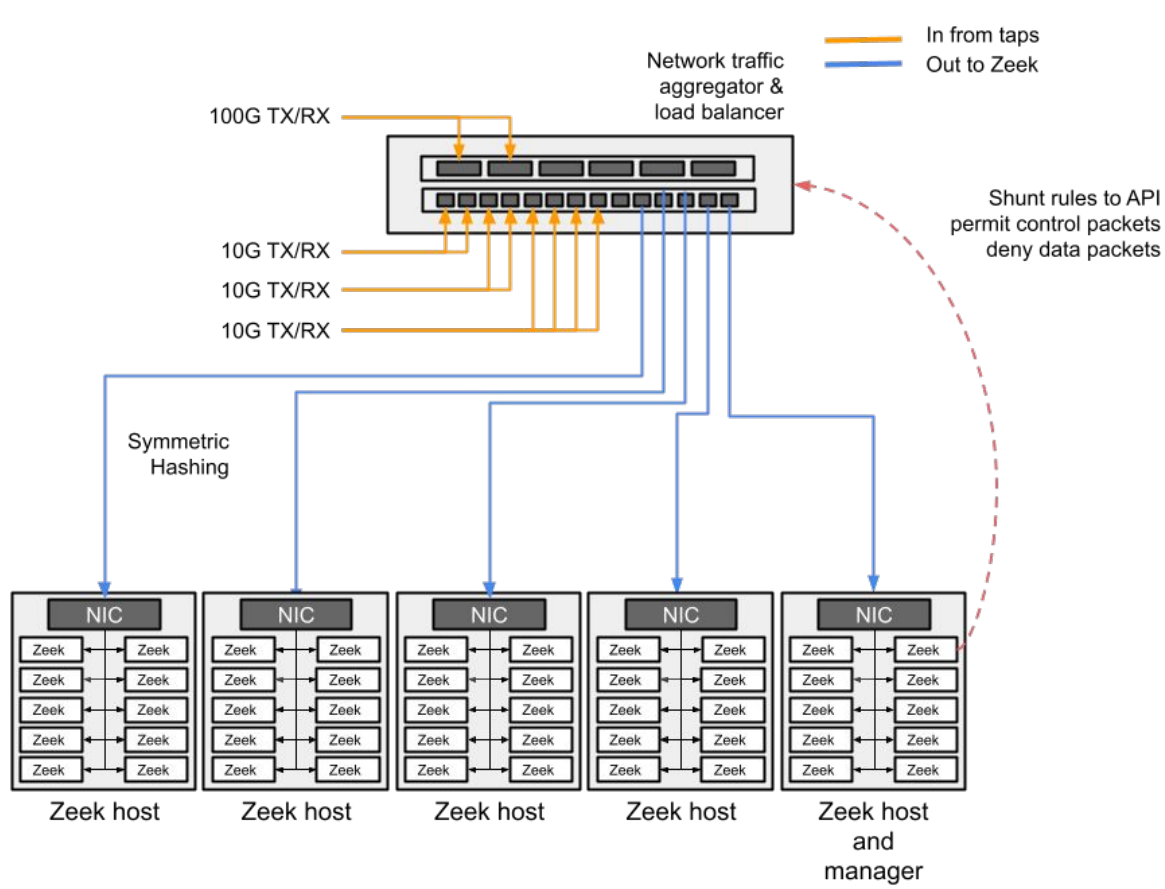


What is a Zeek Cluster?



What / Why

- Zeek processes started with scripts that tune their purpose
 - As opposed to running a single Zeek process, reading from a pcap, for instance
 - Can run on one or more pieces of hardware
- Scales quite nicely
- Supporting tools handle tasks like monitoring processes and log rotation
- If you need 24x7 network monitoring, clustering is your answer



Source : Zeek Project. Creative Commons



Sniffing Packets

- SPAN/MIRROR ports
 - Configured on the router/switch
 - Pros :
 - Flexible (could mirror only certain network segments)
 - Might keep your network people happy (they can control the setup)
 - Cons :
 - Consumes router resources (CPU/physical ports)
 - Vulnerable/Visible to attackers (modifying/viewing configs)
 - Mistakes in managing network gear can impact flows



Sniffing Packets

- Taps
 - No moving parts, no config, invisible to attackers
 - Each link consumes two ports
 - Passive
 - Split light at fixed ratios (50/50 or 70/30 are common)
 - Light budget requires coordination with networking team based on distances
 - Active
 - Electrically powered device
 - Useful for converting signal types (such as 10 Gb SR to 10 Gb LR)
 - Useful for copper infrastructure



Tap Aggregation

- Merges multiple tap sources
- Can output to groups of tools
- Arista / Gigamon / Ixia, for example
- Advanced features
 - Packet deduplication
 - Packet Slicing
 - Timestamping
 - Tagging



Worker Load Balancing

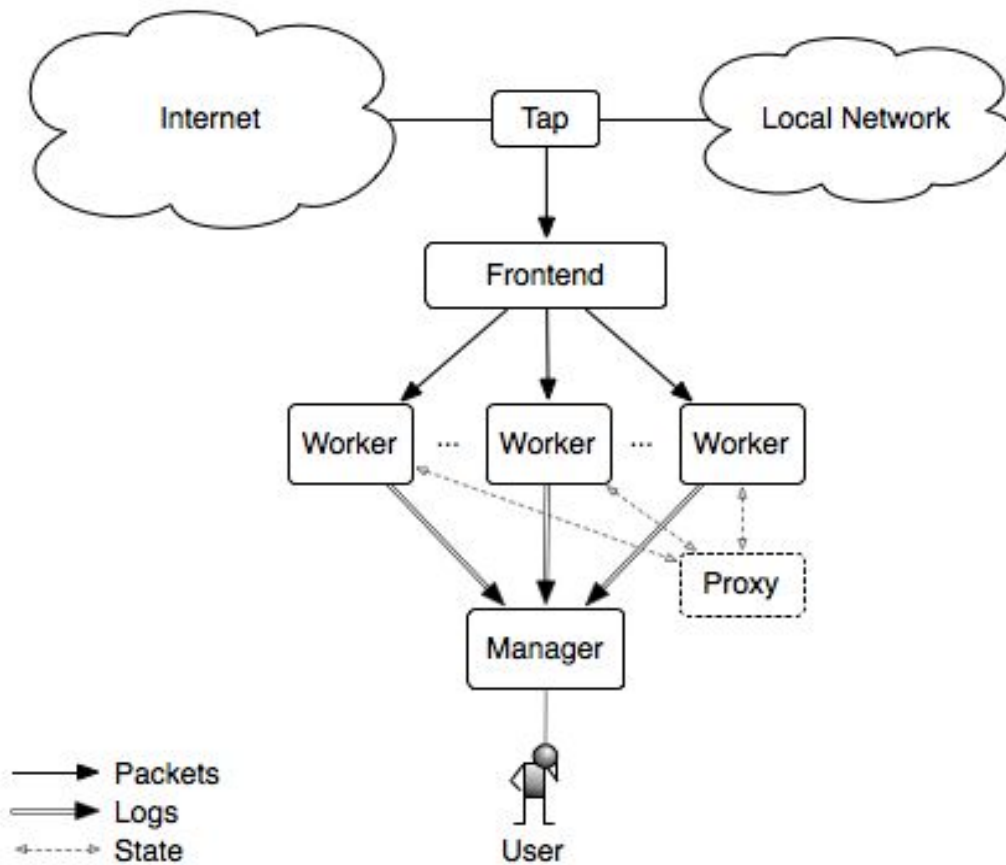
- Commercial Solutions

- Examples : Endace / Napatech / Myricom
- Expensive solution for very high performance
- I'm most familiar with Myricom
 - Provides in-memory ring buffers with symmetric hashing
 - Very easy to support
 - Proprietary driver, updates can lag behind kernel versions



Worker Load Balancing

- Open Source
 - AF_PACKET
 - Built into Linux kernel
 - Can be challenging to setup and tune
 - Interrupt-based
 - PF_Ring
 - External modules, recompile kernel for upgrades
 - Compile zeek with the correct pf_ring libs
 - DPDK
 - [Zeek package](#) handles configuration
 - Also check out Vlad's [talk at ZeekWeek](#)
 - Comes with ~ 50 NIC drivers
 - Polling-based
 - Does not currently support bonded interfaces



Source : Zeek Project. Creative Commons

Cluster Components

- Manager
- Worker
- Proxy
- Logger





Cluster Components

- Manager
 - Writing logs from workers (when not using separate logger nodes)
 - Useful for smaller volumes or limited hardware
 - Can also run scripts when you want to conduct analysis across multiple workers.

Cluster Components



- Manager
- Worker
 - Analyze network packets to create (most) logs
 - Run scripts to alert on interesting traffic



Cluster Components

- Manager
- Worker
- Proxy
 - Data storage (can be distributed with multiple proxies)
 - State sharing among cluster workers
 - Built on Broker (topic-based pub/sub communication using C++ Actor Framework)
 - Start with one, then add more if load is an issue



Cluster Components

- Manager
- Worker
- Proxy
- **Logger**
 - Write logs
 - Easy support to load balance among multiple logger nodes
 - Ideal way to offload log I/O in high volume environments
 - Defaults to TSV format, but easy to switch to JSON

Cluster Architecture



- Single Physical Box
 - Inexpensive, easy management
 - Useful for a single feed



Cluster Architecture

- Hardware Cluster
 - Separate boxes for workers
 - 1U Pizza boxes
 - Fast CPU, Lots of RAM, small disk (low I/O requirements)
 - Allows horizontal scaling
 - Not good candidates for VM environments
 - Manager/Proxy/Logger on a single host
 - Moderate CPU/RAM
 - Fast Disk, ideal for logger node(s)
 - Good components to consider running in a VM environment

Cluster Architecture



- Networking
 - 1Gb interfaces more than adequate for inter-process comms
 - Broker used to share state, make queries, send logs to manager/logger
 - SSH used by Zeekctl (more on that later!)



Performance


- How big?
 - That depends. A lot.
 - ~ 6 GB per worker process
 - Faster CPUs are always better
- LBL 100G Paper
 - Good architecture overviews
- SEPTun-Mark-II
 - Suricata focused, but deep dive on NUMA, OS tuning, etc.
- CPU Pinning
 - Avoid using hyperthreaded cores for workers
- Fine tuning ethtool - tro/tso/rx/tx/ etc..
- Metrics/Monitoring
 - Prometheus - Zeek Exporter from ESNET
 - Telemetry FW
 - Collectd



Shunting

- Reduce load by blocking elephant flows
- Streams/large file transfers over TLS, for example
- Ideally, you block the data packets, but let TCP control packets through
 - Zeek can tell you conn size, packet count, duration
- Dumbno (for Arista)

Shunting



Real World Deployments


- Indiana University
- University of Delaware
- ESnet



N/w Architecture for Zeek deployment # 1




At ESnet

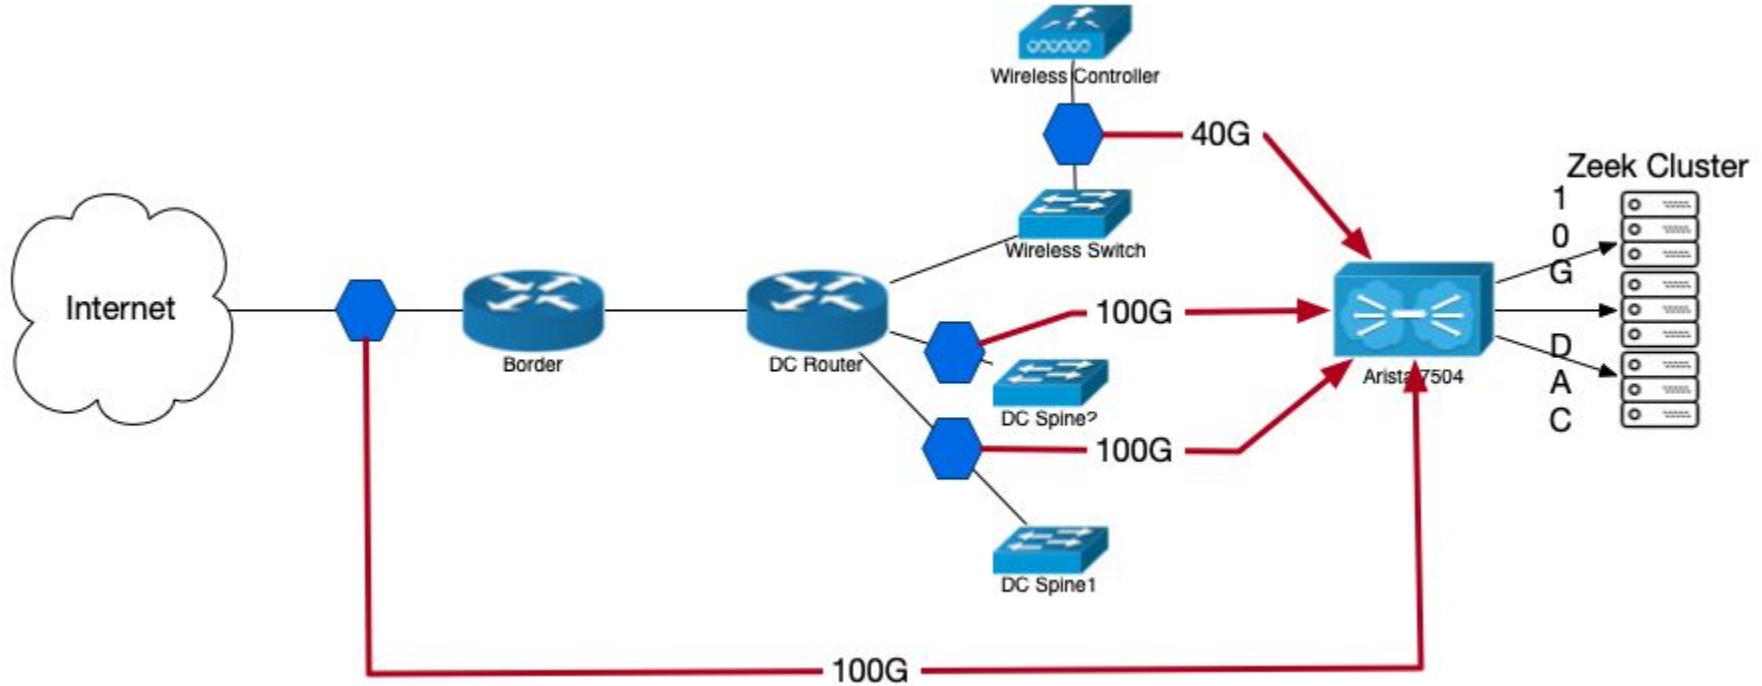


N/w Architecture for Zeek deployment #2

At UD



Indiana University (One Campus Subset)





Filesystem layout



Filesystem Layout

- Not an exhaustive tour!
- Default Path : `/opt/zeek/` or `/usr/local/zeek`
 - `bin/`
 - `etc/`
 - `include/` - ignoring today
 - `lib/` `var/lib` - ignoring today
 - `logs/`
 - `share/`
 - `spool/`



Filesystem Layout

- bin/
 - zeek-cut
 - Extract columns from zeek logs (non-JSON)
 - Convert Unix epoch
 - zeek
 - Zeek executable with various options to run
 - zkg
 - Zeek package manager executable
 - zeekctl
 - Cluster management tool
 - Python script, more about it later



Filesystem Layout

- etc/
 - network.cfg
 - Define your local networks
 - node.cfg
 - Configure a cluster
 - Extra features (lb_procs, pin_cpu, lb_method)
 - zeekctl.cfg
 - Configure cluster management tool
 - Set shorter rotation time



Filesystem Layout

- logs/
 - current
 - Default Log Path (cluster-mode)
 - current directory
 - Default Log Path (CLI mode, logs overwritten each run)
 - YYYY-MM-DD
 - Logs archived to this dir, rotated hourly by default



Filesystem Layout

- spool/
 - state.db
 - cluster worker state
 - zeekctl-config.sh
 - Cluster management variables
 - worker-1
 - .pid
 - .cmdline
 - .env_vars
 - logger
 - logs (current is just a symlink)



Filesystem Layout

- share/
 - zeek/base/
 - Base scripts
 - zeek/policy/
 - Additional scripts
 - zeek/site/
 - Site specific scripts
 - <Zkg packages dir>



Docker / Zeek Usage

Docker Setup

- Github Link- Docker How to doc
 - <https://github.com/zeek/zeek-training/tree/master/Intro-to-Zeek22>
- Docker image
 - Zeek v6.0.1 installed on the image- <https://hub.docker.com/repository/docker/zeek/training>
 - `# docker pull zeek/training:intro23`
 - `# docker run -it --privileged zeek/training:intro23 /bin/bash`
 - Path inside the container to training resources used during this training: `/zeek/training-res/`

Verifying if Zeek installed correctly:

```
# /zeek/bin/zeek -h
```

```
# export PATH=/zeek/bin:$PATH
```



MORNING BREAK [30 mins]



Exercise#1

Fire up a cluster

- Edit node.cfg - comment out the standalone section, and uncomment the logger, proxy, manager and worker-1 sections.
 - `# vi /zeek/etc/node.cfg`

```
//You have to compile Zeek with pf_ring
[worker-1]
type=worker
host=10.0.0.50
interface=eth0
lb_method=pf_ring
lb_procs=10
pin_cpus=2,3,4,5,6,7,8,9,10,11
```

```
//Supported out-of-the-box
[worker-1]
type=worker
host=localhost
interface=af_packet::eth0
lb_method=custom
lb_procs=8
pin_cpus=0,1,2,3,4,5,6,7
```

- Set shorter rotation time
 - `# vi /zeek/etc/zeekctl.cfg`
`LogRotationInterval = 180 //(3mins)`
- Set your local network
 - `# vi /zeek/etc/network.cfg`



- turn off the checksum with Zeek redef

```
# vi /zeek/share/zeek/site/local.zeek
    redef ignore_checksums = T;
```
- Then run:
 - `# zeekctl deploy`
- Make sure cluster is running:
 - `# zeekctl status`
- Create some traffic for the cluster to sniff
 - `# wget google.com`
 - `# ls /zeek/logs/current/`
- Send few pcaps to the sniffing interface:
 - `# tcpreplay -i eth0 --mbps=500 /zeek/training-res/*.pcap`
 - `# ls /zeek/logs/current`
- Stop the cluster
 - `# zeekctl stop`



Customizing Zeek

- `/zeek/share/zeek/site/local.zeek`
 - local customizations
 - Can be done inline or separate script
 - enable or disable scripts
 - change variables



Customizing Zeek

Loading additional scripts:

```
# nano /zeek/share/zeek/site/local.zeek
```

<Changes to local.zeek require restart in cluster mode>

```
@load policy/protocols/http/header-names
```

```
zeek -h  
$ZEEKPATH
```

Zeek file search path order :

- . (current directory)
- /usr/local/zeek/share/zeek
- /usr/local/zeek/share/zeek/policy
- /usr/local/zeek/share/zeek/site



Customizing Zeek

- Zeek scripts have an export section.
 - Redefinable variables
 - main.zeek is best place to start looking

```
# less /zeek/share/zeek/base/protocols/http/main.zeek
```

```
export {
```

```
# This setting changes if passwords used in  
# Basic-Auth are captured or not.  
option default_capture_password = F;
```

```
}
```



Customizing Zeek

Value

```
---  
redef SSH::password_guesses_limit = 15;
```

Table

```
---  
redef SSH::ignore_guessers += {  
    [192.168.1.2/32] = 192.168.1.20/32 };
```

Boolean

```
---  
redef HTTP::log_server_header_names = T;
```



Config Framework

Change values on the fly

```
module Test;
```

```
export {  
  option enable_feature = F;  
}
```

```
@redef Test::enable_feature = T;
```

Define config option file

```
---  
redef Config::config_files += { "/zeek/share/zeek/site/config.dat" };
```

Format

```
---  
[option name][tab/space][new value]
```

```
Test::enable_feature    T
```



Customizing Zeek (10 min)

```
# vi /zeek/share/zeek/site/local.zeek
```

Log server header names in HTTP (in local.zeek)
`@load policy/protocols/http/header-names`

Start with : `/zeek/training-res/http-auth.pcap`

```
# tcpreplay -i eth0 --mbps=500 /zeek/training-res/http-basic-auth.pcap
```

Log http usernames and passwords
`redef HTTP::default_capture_password=T;`

```
# tcpreplay -i eth0 --mbps=500 /zeek/training-res/http-auth.pcap
```



Package Manager (zkg)

- Package Manager
 - zkg
 - <https://packages.zeek.org>
- Can add new scripts or analyzers
 - Analyzer examples :
 - QUIC
 - PerfSonar
 - Scripts :
 - BZAR
 - DHCP OUI data



Let's Use zkg!

- List all zeek packages and install a few
 - # zkg refresh
 - # zkg list all
 - # zkg install cve-2020-0601
- Check if they got installed
 - # zkg list installed
 - # ls /zeek/share/zeek/site
- Let's load the packages
 - uncomment *@load packages* in site/local.zeek



Let's Use zkg!

- zkg load/unload and info
 - # zkg unload cve-2020-0601
 - # zkg info cve-2020-0601
- Remove and uninstall all the packages
 - # zkg purge
 - # zkg list installed
- Other useful zkg commands
 - remove, pin, unpin, upgrade, refresh



Zkg Exercise

10 minutes..

Running Zeek with zkg installed packages, specifically with JA3.

```
# zkg install ja3
# zeekctl deploy
# tcpreplay -i eth0 --mbps=500 /zeek/training-res/extract.pcap
# less /usr/local/zeek/logs/current/ssl.log
# cat ssl.log | zeek-cut -d ja3 ja3s
```



Zkg Exercise

Extract a file from a file transfer connection using file-extract [zkg package](#)

```
# zkg install file-extraction
```

Add all file types extraction to the config of the package:

```
# vi /zeek/share/zeek/site/packages/file-extraction/config.zeek
redef path = "/extracted_files/";
@load ./plugins/extract-all-files.zeek
```

```
# mkdir /extracted_files
# zeekctl deploy
# tcpreplay -i eth0 --mbps=500 /zeek/training-res/extract.pcap
# ls /extracted_files/
# less files.log
```



LUNCH Break
[Be back by 2pm]



HAND OFF TO CHRISTIAN



The Logging Framework



Packets → **Zeek** → **?**



Packets → **Zeek** → **Logs!**



Logging Framework: Basic Purpose

- Manages all aspects of log writes in Zeek
- Provides APIs to Zeek scripts to
 - Define logs and their structure
 - Configure output technologies and their properties
 - Write individual log entries
 - Extend a log's structure
 - Adapt and filter log writes
- Zeek's default scripts provide ~60 logs
- Main reason people start using Zeek!



Logging Framework: Main Concepts

- Streams — what to log
 - Corresponds to a single log (for example, conn.log)
 - Defines fields with name and type
- Filters — how & whether to log
 - Determines what gets written and where
 - Each stream has a default filter that just writes everything
- Writers — where to log to
 - Defines output format
 - Zeek ships with ASCII writer (tab-separated or JSON) and SQLite writer
 - Other available via Zeek packages



Logging Framework: Defining a Log

```
# Create ID for our new stream. This identifies the log in a module.
redef enum Log::ID += { LOG };

# Define (or redef, in order to extend) the record type to log.
# This defines all log columns and data types:
type Info: record {
    ts: time          &log;
    id: conn_id       &log;
    service: string &log &optional;
    missed_bytes: count &log &default=0;
};

event zeek_init() {
    # Create the stream. This adds a default filter that writes ASCII logs.
    Log::create_stream(Foo::LOG, [$columns=Info,
                                   $path="foo"]);
}
```



Logging Framework: Writing a Log Entry

```
# Identify a suitable place to write the log entry.  
# This is usually in an event handler; often when  
# Zeek is done with a connection:  
event connection_state_remove(c: connection) {  
    local rec: Foo::Info = [$ts=network_time(), $id=c$id];  
    Log::write(Foo::LOG, rec);  
}
```



Logging Framework: Behold foo.log!

```
#separator \x09
#set_separator ,
#empty_field (empty)
#unset_field -
#path foo
#open 2019-04-07-00-27-05
#fields ts id.orig_h id.orig_p id.resp_h id.resp_p service missed_bytes
#types time addr port addr port string count
1052146262.950001 203.241.248.20 3051 80.4.124.41 80 - 0
#close 2019-04-07-00-27-05
```



Logging Framework: Behold foo.log (in JSON)

```
{"ts":1052146262.950001,  
  "id.orig_h":"203.241.248.20",  
  "id.orig_p":3051,  
  "id.resp_h":"80.4.124.41",  
  "id.resp_p":80,  
  "missed_bytes":0}
```

- Enabled by an ASCII writer configuration setting:

```
redef LogAscii::use_json = T;
```

- Also see the [add-json](#) Zeek package



Logging Framework: Add a Field to conn.log

```
redef record Conn::Info += {  
    is_private: bool &default=F &log;  
};  
  
event connection_state_remove(c: connection) {  
    if ( c$id$resp_h in Site::private_address_space )  
        c$conn$is_private = T;  
}
```

- Note how this just adds state, nothing changes about the existing log write(s)
- Available in training material: add the following to local.zeek and redeploy.
[@load /zeek/training-res/log-private.zeek](#)
- Replay some traffic to trigger log writes:
[# tcpreplay -i eth0 --mbps=500 /zeek/training-res/mal.pcap](#)



Logging Framework: Disable a Whole Log

```
event zeek_init() {  
    Log::disable_stream(Syslog::LOG);  
}
```




Logging Framework: Steering Output

```
function myfunc(id: Log::ID, path: string, rec: HTTP::Info) : string {  
    local r = Site::is_local_addr(rec$id$resp_h) ? "local" : "remote";  
    return fmt("%s-%s", path, r);  
}
```

```
event zeek_init() {  
    Log::remove_filter(HTTP::Log, "default");  
    local filter: Log::Filter = [$name="http-split",  
$path_func=myfunc];  
    Log::add_filter(HTTP::LOG, filter);  
}
```

- Available in training material: add the following to local.zeek and redeploy.
`@load /zeek/training-res/split-path.zeek`
- Replay some traffic to trigger log writes:
`# tcpreplay -i eth0 --mbps=500 /zeek/training-res/*.pcap`



Logging Framework: Filtering Log Writes

- Built on the concept of hooks in the Zeek language
 - Hooks are like events but synchronous: think callbacks.
 - One hook can have multiple handlers, any of which can shortcut further processing.
- `Log::log_stream_policy(rec: any, id: ID);`
 - Global policy hook, invoked for every log write
 - Great way to interpose on any log write in the system
- `Log::PolicyHook: hook(rec: any, id: ID, filter: Filter);`
 - Each log can also define its own policy hook, for filtering specific to this log.
- Issuing a “`break;`” in a hook will veto the log write.



Logging Framework: Further Reading

- Read the docs: <https://docs.zeek.org/en/master/frameworks/logging.html>
- Watch Justin's [talk at ZeekWeek'22](#).
- Real-world examples of log additions/tweaks:
 - This adds a [Community ID](#) field to conn.log
`@load protocols/conn/community-id-logging`
 - An interesting example of adding columns
<https://github.com/corelight/log-add-vlan-everywhere>
 - Packages that add parsers usually add logs too
<https://github.com/cisagov/icsnpp-bacnet>



The Notice Framework



Notice Framework

- Notices convey interesting activity that Zeek encounters. They can be alerts, depending on your site's policy. Many places in Zeek trigger notices, and you can add more.
- Notices have actions:
 - `ACTION_LOG`, `ACTION_ALARM`, `ACTION_EMAIL`, ...
- Hooks allow applying actions and modification of notices before they're sent on to the action plugins:
 - `hook Notice::policy(n: Notice:Info)`
 - Hooks can have priorities (default is 0, higher means earlier)
 - Hooks can abort later hook bodies with the `break` keyword.

Notice Framework



Raising Notices

- **NOTICE** function is called to raise a notice for any occurrence that a user may want to be notified about or take action on:

```
NOTICE([$note=Password_Guessing,  
        $msg=fmt("%s appears to be guessing SSH passwords (seen in %d conns).",  
        key$host, r$num),  
        $src=key$host,  
        $identifier=cat(key$host)]);
```

Notice Framework



Automated Suppression

- Notices have built-in suppression (to avoid repeated events)
 - Generally controlled/defined by the script writer
 - “intrinsically duplicate” notices suppression is implemented with an optional field in `Notice::Info` records named `$identifier` which is a simple string.
 - `$suppress_for` optional field can be used to configure custom interval.



Notice Exercise

- Install scan detection:
zkg install bro-simple-scan
- Uncomment @load packages in local.zeek
- Run against mal.pcap
zeekctl deploy
tcpreplay -i eth0 --mbps=500 /zeek/training-res/*.pcap
- Look at the generated notice.log file



The Telemetry Framework

Telemetry Framework



- This framework provides metrics about Zeek's internal state, as well as characteristics of the monitored traffic.
- It's fairly Prometheus-inspired and supports scraping.
- Zeek ships with some metrics by default. You can add others by scripting.
- These metrics are not a Zeek log data export mechanism.



Telemetry Framework

- Supports the same metric types as most Prometheus client libraries, with the exception of the Summary type.
- Metrics types:
 - Counter - Continuously increasing, resets on process restart.
 - Gauge - Gauge metric can increase and decrease
 - Histogram - Pre-configured buckets of observations.



Telemetry Framework

In a cluster setting:

- Metrics can be collected on manager node.
 - No aggregation of metrics happen at the collecting node
 - Each node has its own metric registry independent of the other nodes.
 - The centralized metrics receive an additional “endpoint” label that can be used to identify the originating node.



Telemetry Framework

Logging and Publishing of metrics:

- By default telemetry.log is written for the metrics of each node in a cluster.
- To allow scraping, the manager process can be configured to run a HTTP server on port 9911/tcp for Prometheus exposition:

```
@load frameworks/telemetry/prometheus
```

- To just scrape available metrics without a cluster setup, try:

```
$ zeek exit_only_after_terminate=T Broker::metrics_port=9911/tcp
```



Telemetry Framework - Exercise

- Load the following script in local.zeek
vi /zeek/share/zeek/site/local.zeek

@load frameworks/telemetry/prometheus
- Redeploy the zeek cluster
zeekctl stop

zeekctl deploy
- Scrape the metrics exposed on port 9911
curl -s localhost:9911/metrics | grep version



Telemetry Framework - Exercise cont.

- Scrape the metrics exposed on port 9911

```
# curl -s localhost:9911/metrics | grep manager
```

```
# curl -s localhost:9911/metrics | grep log_writes
```
- Look at the default telemetry.log file

```
# less /zeek/logs/current/telemetry.log
```

```
# less /zeek/logs/current/telemetry.log | grep manager
```

Telemetry Framework - Custom Metrics Example



- Look at the custom script to count the number of http requests:
cat /zeek/training-res/http-counter.zeek



Telemetry Framework - Custom Metrics Example

- Load the following script in local.zeek
#vi /zeek/share/zeek/site/local.zeek

@load /zeek/training-res/http-counter.zeek
- Redeploy the zeek cluster
zeekctl stop

zeekctl deploy
- Scrape the metrics exposed on port 9911
curl -s localhost:9911/metrics | grep monitored_http_requests



AFTERNOON BREAK [30 mins]



Meet the

Management Framework

[policy/frameworks/management](#)



Goals of this effort

- Establish modern cluster management for Zeek
 - zeekctl's model isn't a great fit for today's infrastructure
 - ssh, rsync, and command invocations
 - Lots of Zeek-external activity — but also solid & featureful!
- Zeek-based Service model
 - Zeek as a service manager
 - Events for messaging
 - Build on Zeek's supervisor framework
 - Be container-friendly
- Feature parity with zeekctl is not a goal, but a guideline

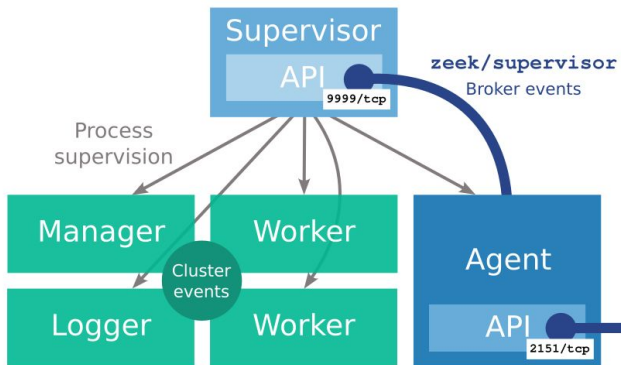


Expectation management

- The Management Framework is an emerging feature
 - Great for single-machine settings & exploring clusterized Zeek
 - Multi-machine settings not fully supported yet
- For production settings, zeekctl remains the right choice
- We plan to deprecate zeekctl in the Zeek 7 timeframe


Cluster System

```
# zeek -j policy/frameworks/management/agent
```



Controller System

```
# zeek -j policy/frameworks/management/controller
```



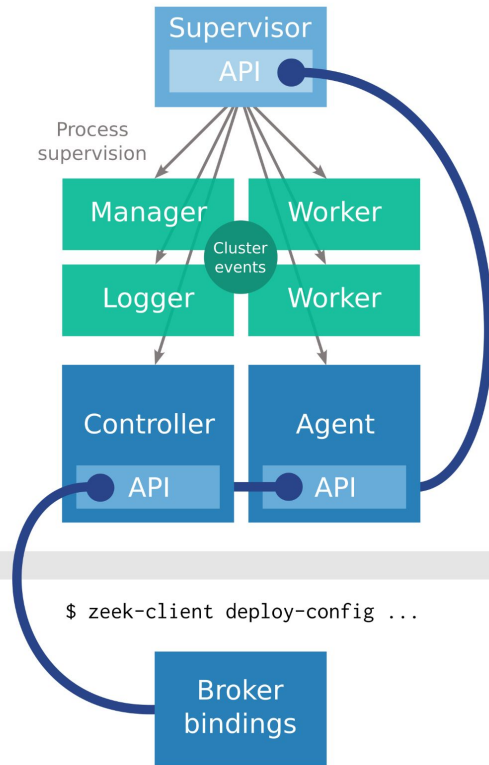
Admin System

```
$ zeek-client deploy-config ...
```



All-in-one System

```
# zeek -j policy/frameworks/management/agent \
policy/frameworks/management/controller
```





Check out the documentation:

<https://docs.zeek.org/en/master/frameworks/management.html>



Let's get started!

- Prepare two shells in a fresh container

To launch additional shells in the existing container, try:

```
$ docker exec -it <container> bash
```



- Launch an all-in-one system

```
# zeek -j policy/frameworks/management/controller  
policy/frameworks/management/agent
```

Note the “-j”: we’ll use Zeek’s Supervisor, a built-in process manager

- Inspect the result

```
# netstat -tplt  
# pstree -pTU 0
```



- Meet zeek-client, the default client CLI

```
# zeek-client --help
# zeek-client show-config
# zeek-client --set <TAB>
```

- Inspect available management instances

```
# zeek-client get-instances
```

zeek-client's output is generally in JSON, suitable for piping into jq etc.

- Write a first cluster config: cluster-minimal.cfg



```
[manager]
role = manager
```

```
[logger]
role = logger
```

```
[worker-01]
role = worker
interface = lo
```



- Staging and retrieving the configuration

```
# zeek-client stage-config cluster-minimal.cfg  
# zeek-client get-config  
# zeek-client get-config --as-json
```

- Ctrl-C the Zeek tree, start it up again, retrieve again

```
# zeek-client get-config
```

The controller has persisted the configuration, reloaded it, and is serving it again.



- Let's run it!

```
# zeek-client deploy
```

- Retrieve current node state

```
# zeek-client get-nodes
```

A view across all instances showing Zeek nodes with their roles, ports, and PIDs.

- Ctrl-C the Zeek tree, start it up again, retrieve nodes again

```
# zeek-client get-nodes
```

The controller has reloaded and re-deployed the configuration.



- Let's make a mistake

Edit the configuration: add another worker, give it an invalid interface.
(See cluster-error.cfg)

- Upload and deploy this, in a single command

```
# zeek-client deploy-config cluster-error.cfg
```



- Let's tweak logging and run some traffic

To study logging, we'll dial down the log rotation interval. Create `local.zeek` with this line:

```
redef Log::default_rotation_interval = 20 sec;
```

Restart everything with this file included — it propagates through to all nodes:

```
# zeek -j policy/frameworks/management/controller  
        policy/frameworks/management/agent local.zeek
```

- Replay a pcap

```
# tcpreplay -i eth0 -mbps=500 /zeek/training-res/extract.pcap
```

- Let's take a look at log production now.



- Step 1: logs get written in local output folders

Take some time to poke around `${ZEEKROOT}/var/lib/nodes/<node>/`.

- Step 2: log rotation

Rotated logs get placed into `${ZEEKROOT}/spool/log-queue/`.

Log rotation interval: `Log::default_rotation_interval`, as usual, default: 1h.

- Step 3: log archival

Log rotation concludes with log archival into `${ZEEKROOT}/logs/`.

By default, log archival happens once per rotation interval. To adjust, redefine

`Management::Agent::archive_interval`

For details and additional customization knobs, see:

<https://docs.zeek.org/en/master/frameworks/management.html#log-management>



- Restart nodes

```
# zeek-client restart [<node1> ...]
```

Compare affected PIDs in subsequent `get-nodes` outputs.

- Inspect cluster node state

You can retrieve global ID values from any set of Zeek nodes. They get rendered in JSON. A simple one:

```
# zeek-client get-id-value ignore_checksums
```

A more complex one:

```
# zeek-client get-id-value Log::active_streams manager
```



- Multiple instances

Let's approximate typical future deployment: multiple Zeek hosts, each with their own instance, and a separate controller. You need three shells, one for each of the following:

```
# zeek -j policy/frameworks/management/controller
# zeek -j policy/frameworks/management/agent
    /zeek/training-res/management/local.inst1.zeek
# zeek -C -j policy/frameworks/management/agent
    /zeek/training-res/management/local.inst2.zeek
```

Deploy a cluster across these instances:

```
# zeek-client deploy-config cluster-multi-inst.cfg
```

Now explore via `get-nodes`, check out `pstree`, etc.



Adjacent technologies: JavaScript

- Starting with Zeek 6, you can script Zeek in JavaScript as well.
- This is an experimental extension, not a full substitute for the Zeek language.
- This runs a Node interpreter and can use packages from the Node ecosystem.
- Really good for quick prototyping, API interaction with other systems, etc
- Implemented as a Zeek plugin!
- See <https://docs.zeek.org/en/master/scripting/javascript.html> for details.



Adjacent technologies: Spicy

- Spicy is a parser generator for structured data
- Domain-specific language to capture syntax and semantics
- Compiles to a safe C++ runtime that is independent of Zeek
 - Prototypes exist e.g. for Wireshark
- Zeek supports Spicy parsers via a glue layer to its event loop
- For an example, see Zeek's [QUIC parser sources](#) (new with 6.1)
- See <https://docs.zeek.org/projects/spicy/en/latest/>



More Zeek Resources

Zeek Documentation:

<https://docs.zeek.org/en/master/>

Zeek packages:

<https://packages.zeek.org/packages>

Zeek btest pcaps:

<https://github.com/zeek/zeek/tree/master/testing/btest/Traces>

Zeek Github source:

<https://github.com/zeek/zeek>

Zeek online:

<https://try.zeek.org/#/?example=hello>



Thanks for attending!
Enjoy Zeeking

Stay Connected



Website - <https://zeek.org>

Slack - <https://zeek.org/slack>

Discourse - <https://community.zeek.org>

Find out more ways to connect at: <https://zeek.org/community/>